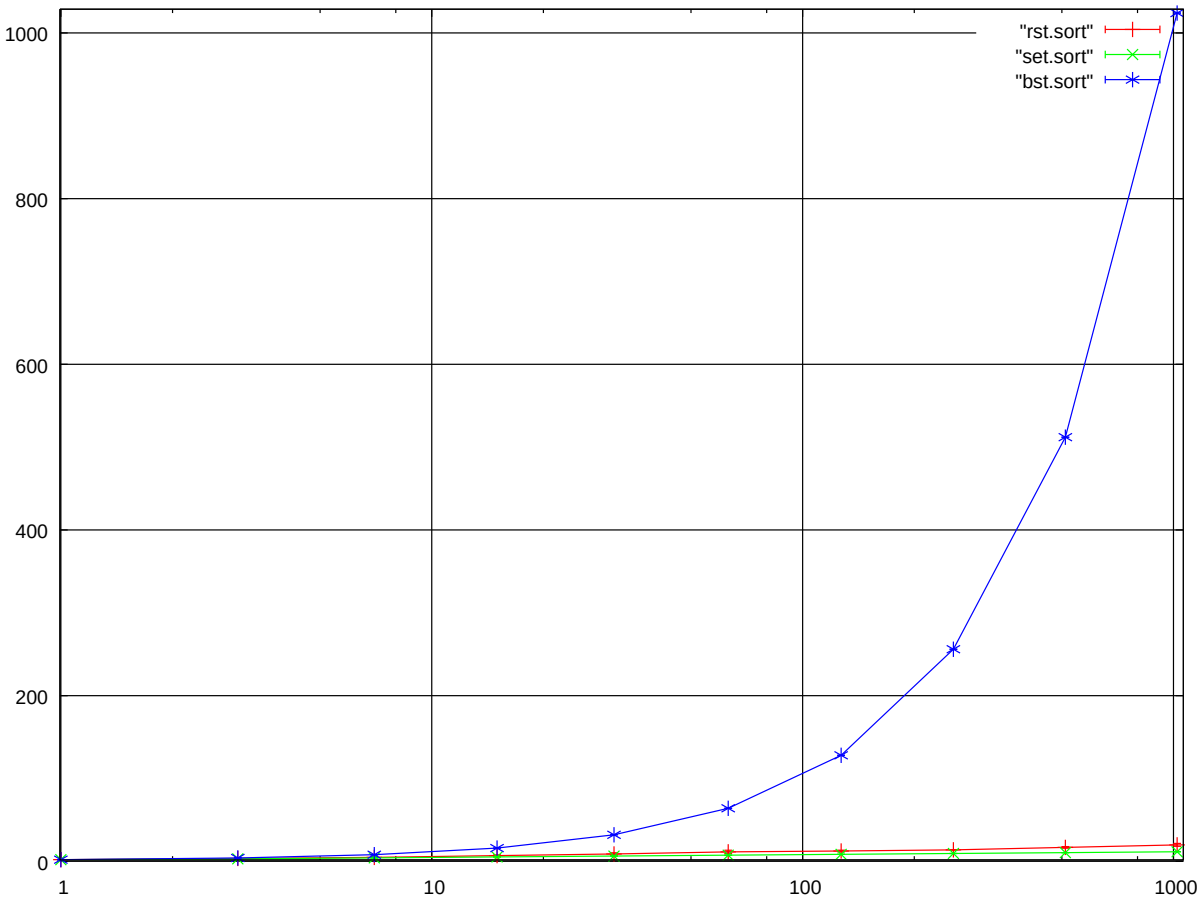

When plotted on a logarithmic scale, the lines for the RST and std::set show a linear relationship between the number of nodes and the average number of comparisons. This agrees with their theoretical big-O of log(N).

On the other hand, the line for the BST is not proportional to log(N). Its line shows exponential growth when graphed on a logarithmic scale.

Both the RST and the RB tree (std::set) greatly outperform the BST in terms of comparisons needed for a successful find operation. The data supports Aragon and Seidel's claim that RSTs are almost as efficient as a RB tree.